

1. Rolling statistics/quantities

Past-window aggregates (mean/std/min/max/percentiles) over several horizons (e.g., 7/30/90 days).

Captures level & variability trends without using the label; quantiles are robust to rare spikes and class imbalance.

On segment-time series (e.g., `Acc_*_RMS`, `Twist`, `Speed`) compute 7/30/90-day stats **excluding the current timestamp** to avoid leakage.

Pros

- Simple, fast, strong baseline
- Quantiles resist outliers (common in rare-event data)

Cons

- Window choice is manual
- May dilute short bursts if windows are too long

Sample :

```
WITH base AS (
    SELECT
        BaseCode AS seg,
        CAST(RecordingDate AS timestamp) AS ts,
        Acc_Vert_RMS, Acc_Lat_RMS, Twist, Speed
    FROM predictive_maintenance.wagondata
)
SELECT
    seg, ts,
    -- 7-day rolling mean/std (past only)
    avg(Acc_Vert_RMS) OVER w7 AS accv_mean_7d,
    stddev_pop(Acc_Vert_RMS) OVER w7 AS accv_std_7d,
    -- 30-day 95th percentile
    percentile_cont(0.95) WITHIN GROUP (ORDER BY Acc_Vert_RMS) OVER w30 AS
    accv_p95_30d,
    avg(Twist) OVER w30 AS twist_mean_30d
FROM base
WINDOW
    w7 AS (PARTITION BY seg ORDER BY ts
              RANGE BETWEEN INTERVAL 7 DAYS PRECEDING AND INTERVAL 1 SECOND
PRECEDING),
    w30 AS (PARTITION BY seg ORDER BY ts
```

```
RANGE BETWEEN INTERVAL 30 DAYS PRECEDING AND INTERVAL 1 SECOND  
PRECEDING);
```

2. Lags & differences

Lagged values, first differences, percent changes, linear trend (windowed slope).

Rare failures often follow **acceleration of degradation**; deltas & slopes reveal momentum.

Per segment, build `lag_1/7/30`, `diff_1`, and `30-day` slope.

Pros

- Highlights deterioration speed
- Low risk of leakage if you only use past lags

Cons

- Sensitive to missingness/irregular sampling

Sample :

```
from pyspark.sql import functions as F, Window as W

w = W.partitionBy("seg").orderBy("ts")

df2 = df \  
    .withColumn("accv_lag1", F.lag("Acc_Vert_RMS", 1).over(w)) \  
    .withColumn("accv_diff1", F.col("Acc_Vert_RMS") - F.col("accv_lag1")) \  
    .withColumn("t_unix", F.col("ts").cast("long")) \  
    # rolling slope over last 30 records (replace with time-based in SQL if needed)  
    .withColumn("accv_slope_30",  
        F.regr_slope(F.col("Acc_Vert_RMS"), F.col("t_unix")).over(w.rowsBetween(-30, -1)))  
    # NOTE: use SQL time windows when you need "last 30 days" instead of "last 30 rows".
```

3. Recency-weighted features (EWMA / EWMVar)

Exponentially-weighted moving averages/variances, so the recent past dominates.

In rare-event prediction, very recent anomalies matter more than old ones.

Build EWMA/EWMVar for key sensors ($\alpha \in [0.05, 0.3]$) **per segment**.

Pros

- Emphasizes fresh signals
- One parameter (α) to tune

Cons

- Requires custom logic in Spark (no built-in EWM in SQL)

Sample :

```
import pandas as pd
from pyspark.sql import functions as F, Window as W
from pyspark.sql.types import StructType, StructField, TimestampType,
StringType, DoubleType

@F.pandas_udf("seg string, ts timestamp, accv_ewma double, accv_ewmvar
double")
def ewm_udf(pdf: pd.DataFrame) -> pd.DataFrame:
    # Input is sorted per partition (seg)
    pdf = pdf.sort_values("ts")
    x = pdf["Acc_Vert_RMS"].astype(float)
    # Exponentially weighted mean/var
    ewma = x.ewm(alpha=0.2, adjust=False).mean()
    ewmvar = x.ewm(alpha=0.2, adjust=False).var(bias=False)
    return pd.DataFrame({"seg": pdf["seg"], "ts": pdf["ts"],
                         "accv_ewma": ewma, "accv_ewmvar": ewmvar})

df_ew = df.groupBy("seg").applyInPandas(ewm_udf)
# NOTE: This keeps causality since it scans forward in time per segment.
```

4. Eventization for rare patterns: counts & time-since-last

Convert continuous signals into **event flags** (e.g., above segment-level P95/P99) and engineer:

- event count / share in last 7/30/90 days
- Time-since-last-event

For imbalanced targets, “how often and how recently things went bad” is highly predictive.

Define flags for **Acc_*_RMS**, **Twist**, **LFI = Speed² × |Curvature|**, etc.

Pros

- Directly addresses rarity and recency
- Robust to scale/unit issues

Cons

- Needs good thresholds (use robust, per-segment quantiles)

Sample :

```

WITH b AS (
  SELECT
    BaseCode AS seg,
    CAST(RecordingDate AS date) AS d,
    Acc_Vert_RMS,
    -- Event if value exceeds per-segment 95th percentile (computed
    historically)
    CASE WHEN Acc_Vert_RMS >
      PERCENTILE(Acc_Vert_RMS, 0.95) OVER (PARTITION BY BaseCode)
    THEN 1 ELSE 0 END AS ev
  FROM predictive_maintenance.wagondata
),
daily AS (
  SELECT seg, d, SUM(ev) AS ev_cnt, COUNT(*) AS n
  FROM b GROUP BY 1,2
)
SELECT
  seg, d,
  -- Past-30-day event count (exclude today)
  SUM(ev_cnt) OVER (PARTITION BY seg ORDER BY d
    RANGE BETWEEN INTERVAL 30 DAYS PRECEDING AND INTERVAL 1 DAY
  PRECEDING) AS ev_cnt_30d,
  -- Time since last event in days
  DATEDIFF(d, MAX(CASE WHEN ev_cnt>0 THEN d END)
    OVER (PARTITION BY seg ORDER BY d
      RANGE BETWEEN INTERVAL 365 DAYS PRECEDING AND INTERVAL
      1 DAY PRECEDING))
    AS days_since_last_ev
  FROM daily;

```

5. Exposure & missingness features

Track **how much data** exists in a window: number of observations, share of missing, sensor uptime.

Imbalanced datasets often have **uneven coverage**; low exposure can confound the model (false “safe” sections).

Apply:

For every window, add `obs_count`, `%missing`, `%imputed`.

Pros

- Improves reliability and calibration
- Helps the model “know what it doesn’t know”

Cons

- Correlated with operating schedules (interpret carefully)

Sample :

```
from pyspark.sql import functions as F, Window as W
w = W.partitionBy("seg").orderBy("ts").rowsBetween(-1000, -1) # last
~1000 rows as a proxy window

df_exp = df \
    .withColumn("obs_cnt_win", F.count(F.lit(1)).over(w)) \
    .withColumn("miss_accv_win",
                F.avg(F.when(F.col("Acc_Vert_RMS").isNull(),
1).otherwise(0)).over(w)) \
    .withColumn("miss_twist_win",
                F.avg(F.when(F.col("Twist").isNull(),
1).otherwise(0)).over(w))
# NOTE: Replace row-based windows with time-based windows in SQL when
required.
```

6. Fourier seasonal terms

Encode seasonality with sine/cosine of day-of-year / week-of-year; month dummies if needed.

Temperature cycles (thermal stress) can correlate with breaks; cyclical encoding avoids artificial jumps at year boundaries.

Add `doy_sin`, `doy_cos`, `woy_sin`, `woy_cos`.

Pros

- Cheap signal; works even without explicit weather data

- Smooth, differentiable (good for linear/NN models)

Cons

- Coarse proxy; add real weather later if available

Sample :

```
SELECT
  seg, ts,
  SIN(2*PI() * dayofyear(ts)/365.0) AS doy_sin,
  COS(2*PI() * dayofyear(ts)/365.0) AS doy_cos,
  SIN(2*PI() * weekofyear(ts)/52.0) AS woy_sin,
  COS(2*PI() * weekofyear(ts)/52.0) AS woy_cos
FROM your_time_series;
```

7. Target encoding (leakage-safe)

For each segment (or segment×line), compute the **historical event rate up to t-1** and use it as a prior/propensity feature.

Imbalanced labels benefit from a **smoothed prior**; expanding mean respects time and reduces variance on rare positives.

Join **trainingcontext** (labels) and compute expanding **avg(target)** **excluding current**; use Bayesian smoothing if data is sparse.

Pros

- Strong priority for rare events
- Encodes persistent “hotness” of a segment

Cons

- Must guard against leakage strictly
- Sparse segments need smoothing

Sample :

```
WITH y AS (
  SELECT BaseCode AS seg, CAST(r_date AS timestamp) AS ts, target
  FROM predictive_maintenance.trainingcontext
),
feat AS (
```

```

SELECT
    seg, ts,
    -- Expanding mean up to t-1
    AVG(target) OVER (
        PARTITION BY seg ORDER BY ts
        ROWS BETWEEN UNBOUNDED PRECEDING AND 1 PRECEDING
    ) AS seg_prior_rate
    FROM y
)
SELECT * FROM feat;
-- NOTE: For smoothing: seg_prior_smoothed = (sum(y) + a) / (count + a + b)

```

Reference

Box, GEP 2016, *Time series analysis : forecasting and control*, Fifth edition., Wiley, Hoboken, New Jersey.

Bergmeir, C & Benítez, JM 2012, ‘On the use of cross-validation for time series predictor evaluation’, *Information Sciences*, vol. 191, pp. 192–213.

Coles, S 2001, *An Introduction to Statistical Modeling of Extreme Values*, 1st ed. 2001., Springer London, London.

Che, Z, Purushotham, S, Cho, K, Sontag, D & Liu, Y 2018, ‘Recurrent Neural Networks for Multivariate Time Series with Missing Values’, *Scientific Reports*, 6085, vol. 8, no. 1.

Jones, LA, Champ, CW & Rigdon, SE 2001, ‘The Performance of Exponentially Weighted Moving Average Charts With Estimated Parameters’, *Technometrics*, vol. 43, no. 2, pp. 156–167.